

Week11

EX0. CPU Scheduling

Process	Estimated CPU Cost	Arrives	Priority
A	4	1	1
B	1	2	2
C	3	5	3
D	2	4	4

Time	HRRN	FIFO/FCFS	RR	SJF	Priority
1	A	A	A	A	A
2	A	A	B	A	A
3	A	A	A	A	A
4	A	A	D	A	A
5	B	B	C	B	B
6	D	D	A	D	C
7	D	D	D	D	C
8	C	C	C	C	C
9	C	C	A	C	D
10	C	C	C	C	D
Avg. Turn-around Time	$(4+4+4+8) / 4 = 5$	$(4+4+4+8) / 4 = 5$	$(9+1+4+8) / 4 = 5.5$	$(4+4+4+8) / 4 = 5$	$(4+4+4+7) / 4 = 3.75$

EX1

The design id is quit simple, just follow the way to do the syscall and finally change the attribute priority.

```
static int
user_main(void *arg) {
#ifdef TEST
    KERNEL_EXECVE2(TEST, TESTSTART, TESTSIZE);
#else
    KERNEL_EXECVE(ex1);
#endif
    panic("user_main execve failed.\n");
}
```

Change to run ex1

```
void
set_priority(int priority){
    sys_setpriority(priority);
}
```

```
static int
sys_setpriority(uint64_t arg[]) {
    int priority = (int)arg[0];
    set_priority(priority);
    return 0;
}

static int sys_gettime(uint64_t arg[]){
    return (int)ticks*10;
}



static int (*syscalls[])(uint64_t arg[]) = {
    [SYS_exit] sys_exit,
    [SYS_fork] sys_fork,
    [SYS_wait] sys_wait,
    [SYS_exec] sys_exec,
    [SYS_yield] sys_yield,
    [SYS_kill] sys_kill,
    [SYS_getpid] sys_getpid,
    [SYS_putc] sys_putc,
    [SYS_gettime] sys_gettime,
    [SYS_labschedule_set_priority] sys_setpriority,
};
```

Add syscall

```
int sys_setpriority(int priority){

    return syscall(SYS_labschedule_set_priority,priority);
}
```

change priority

```
 set_priority(int priority)
{
    cprintf("set priority to %d\n",priority);
    current->labschedule_priority = priority;
}

```

The change of .h file is trivial and I think it is not necessary to show.

OpenSBI v0.6

OpenSBI

Platform Name : QEMU Virt Machine
Platform HART Features : RV64ACDFIMSU
Platform Max HARTs : 8
Current Hart : 0
Firmware Base : 0x80000000
Firmware Size : 120 KB
Runtime SBI Version : 0.2

MIDELEG : 0x0000000000000222
MEDELEG : 0x000000000000b109
PMP0 : 0x0000000080000000-0x000000008001ffff (A)
PMP1 : 0x0000000000000000-0xffffffffffffffff (A,R,W,X)
OS is loading ...

memory management: default_pmm_manager
physcial memory map:
memory: 0x08800000, [0x80200000, 0x885fffff].
sched class: RR_scheduler
SWAP: manager = fifo swap manager
setup timer interrupts
The next proc is pid:1
The next proc is pid:2
kernel_execve: pid = 2, name = "ex1".
Breakpoint

-----ex1---start-----
set priority to 5
-----ex1---end-----

The next proc is pid:1
all user-mode processes have quit.
The end of init_main
kernel panic at kern/process/proc.c:413:
initproc exit.

wangxiequn11910405@ubuntu-linux-20-04-desktop:~/Desktop/Parallels Shared Folders
/Home/CLionProjects/week11\$

EX2


```
kernel_execve: pid = 2, name = "ex3".
Breakpoint
main: fork ok, now need to wait pids.
0 The next proc is pid:7
  set good to 2
  The next proc is pid:3
  set good to 3
  The next proc is pid:4
  set good to 1
  The next proc is pid:5
  set good to 4
  The next proc is pid:6
  set good to 5
  child pid 6, acc 4000001
  The next proc is pid:2
  The next proc is pid:5
  set good to 4
  child pid 5, acc 4000001
  The next proc is pid:2
  The next proc is pid:3
  set good to 3
  child pid 3, acc 4000001
  The next proc is pid:2
  The next proc is pid:7
  child pid 7, acc 4000001
  The next proc is pid:2
  The next proc is pid:4
  child pid 4, acc 4000001
  The next proc is pid:2
main: wait pids over
The next proc is pid:1
all user-mode processes have quit.
The end of init_main
kernel panic at kern/process/proc.c:413:
  initproc exit.
```

```
set_good(int good)
{
    cprintf("set good to %d\n", good);
    current->labschedule_good = good;
}
```

```

void
set_good(int good)
{
    cprintf("set good to %d\n", good);
    current->labschedule_good = good;
}

```

```

static struct proc_struct *
alloc_proc(void) {
    struct proc_struct *proc = kmalloc(sizeof(struct proc_struct));
    if (proc != NULL) {
        proc->state = PROC_UNINIT;
        proc->pid = -1;
        proc->runs = 0;
        proc->kstack = 0;
        proc->need_resched = 0;
        proc->parent = NULL;
        proc->mm = NULL;
        memset(&(proc->context), 0, sizeof(struct context));
        proc->tf = NULL;
        proc->cr3 = boot_cr3;
        proc->flags = 0;
        memset(proc->name, 0, PROC_NAME_LEN);
        proc->wait_state = 0;
        proc->cptr = proc->optr = proc->yptr = NULL;
        proc->time_slice = 0;
        proc->labschedule_priority = 1;
        proc->labschedule_good = 6;
    }
}

```

```

static int (*syscalls[])(uint64_t arg[]) = {
    [SYS_exit]          sys_exit,
    [SYS_fork]          sys_fork,
    [SYS_wait]          sys_wait,
    [SYS_exec]          sys_exec,
    [SYS_yield]         sys_yield,
    [SYS_kill]          sys_kill,
    [SYS_getpid]        sys_getpid,
    [SYS_putc]          sys_putc,
    [SYS_gettime]       sys_gettime,
    [SYS_labschedule_set_priority] sys_setpriority,
    [SYS_labschedule_set_good] sys_setgood
};

```

```

#define SYS_labschedule_set_priority 255
#define SYS_labschedule_set_good 256
/* SYS fork flags */

```

The good is like a priority to set. For init, set to 6

We close the clock and change pick function.

```

static struct proc_struct *
RR_pick_next(struct run_queue *rq) {
    list_entry_t *le = list_next(&(rq->run_list));
    list_entry_t *cu = list_prev(le);
    struct proc_struct * tp = le2proc(le, run_link);
    struct proc_struct * cp = le2proc(cu, run_link);
    struct proc_struct * max_p = cp;

    while (tp!=cp){
        if(max_p->labschedule_good<tp->labschedule_good){
            max_p = tp;
        }
        le = list_next(le);
        tp = le2proc(le, run_link);
    }

    return max_p;
}

```

For every pick, choose the biggest value of good.